

Kubernetes

tuna



Kubernetes

Open Source

<https://github.com/kubernetes/kubernetes>

Very active open source project

12k stars, 750+ contributors

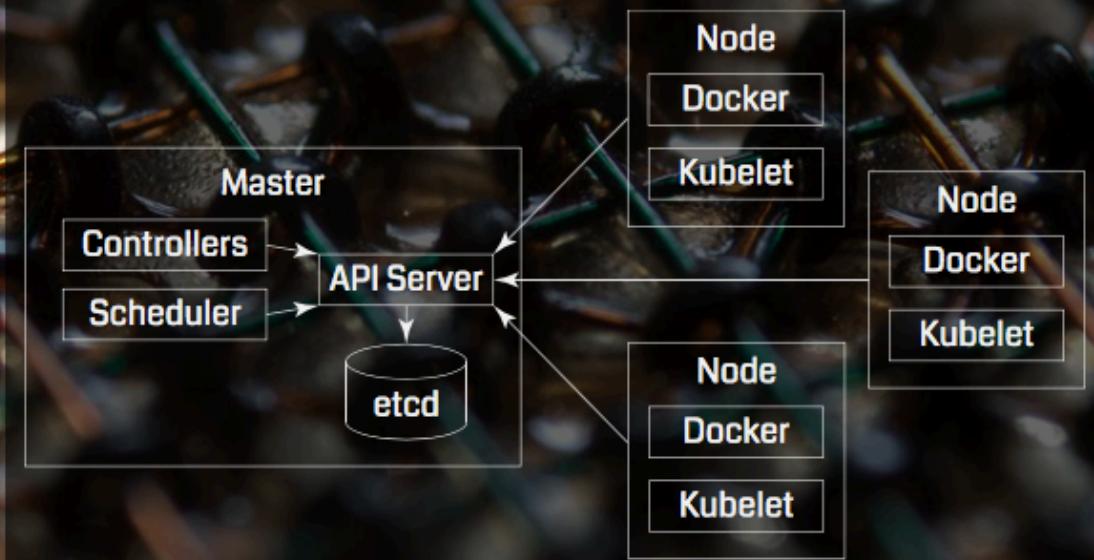
Apache 2 licensed

Written in Go

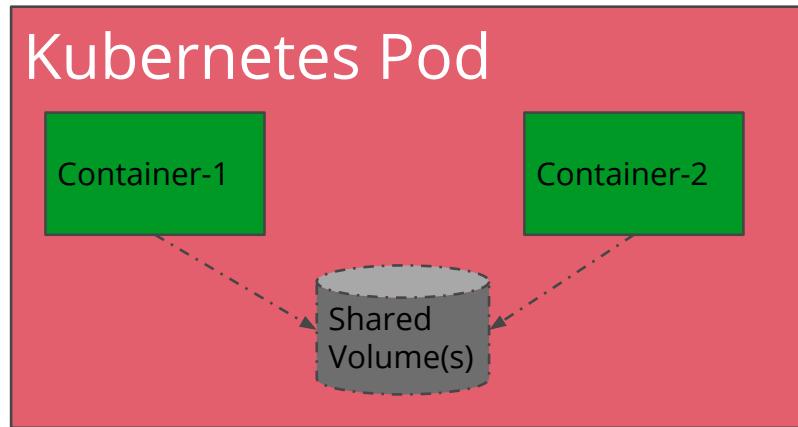
In progress: donate to the Cloud Native Computing Foundation

Core Concepts

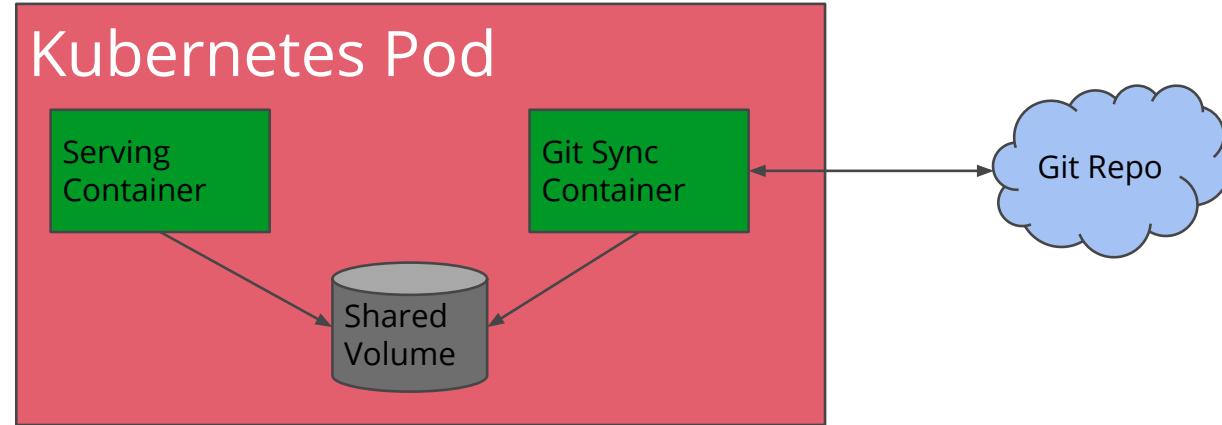
Cluster



Components: Pods

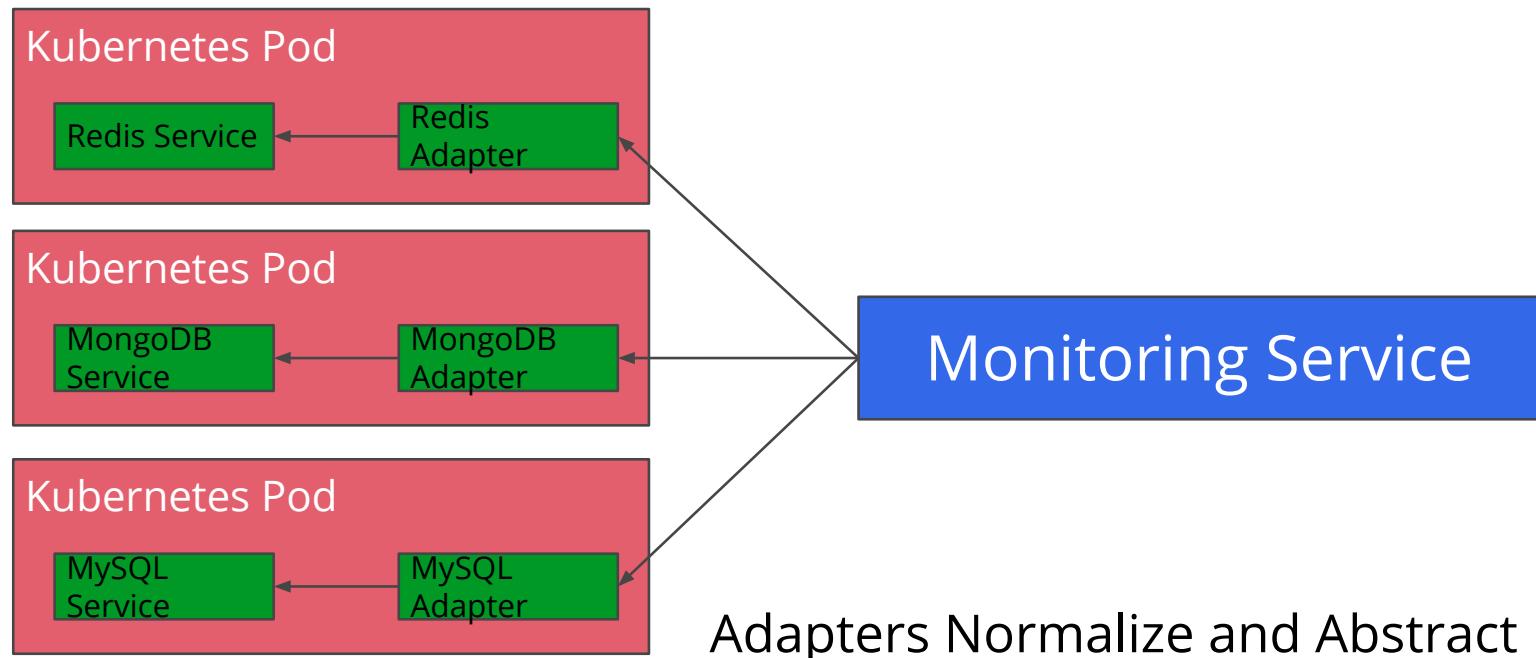


Components: Sidecars

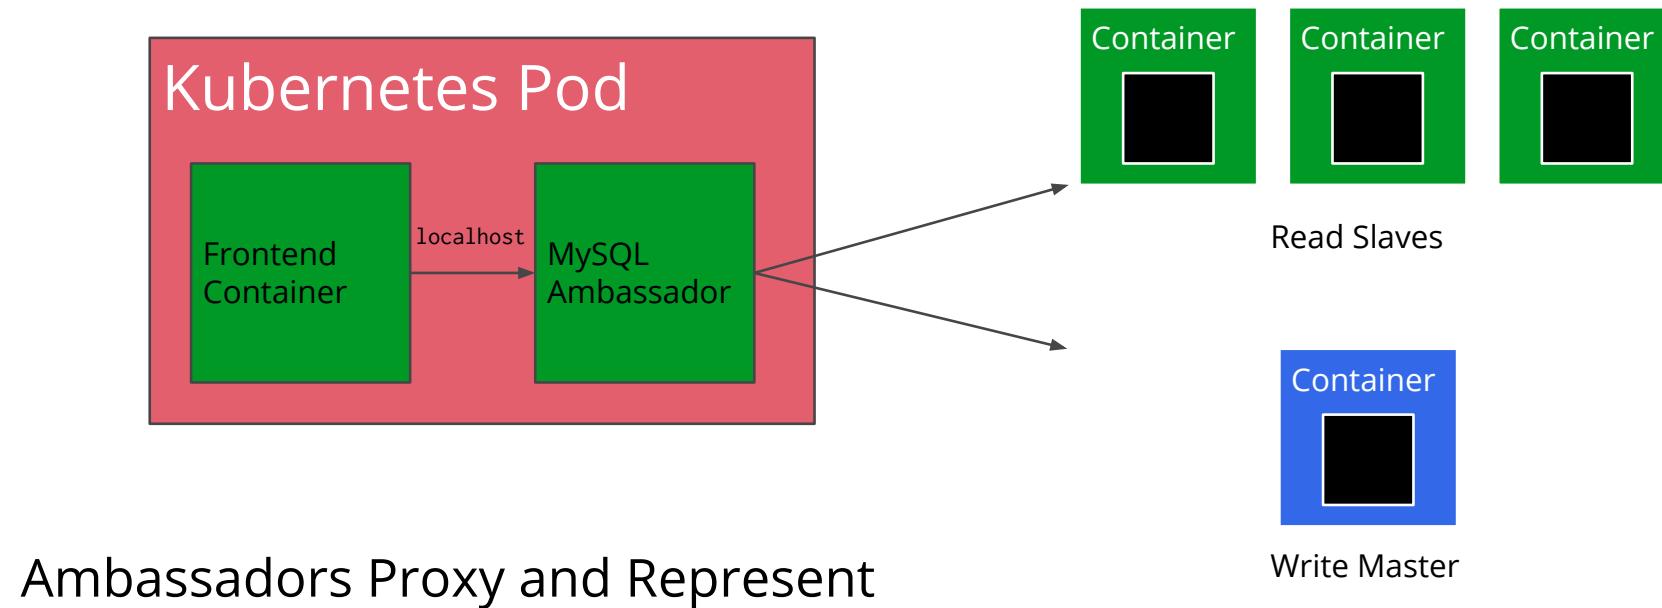


Sidecars Extend and Enhance

Components: Adapters

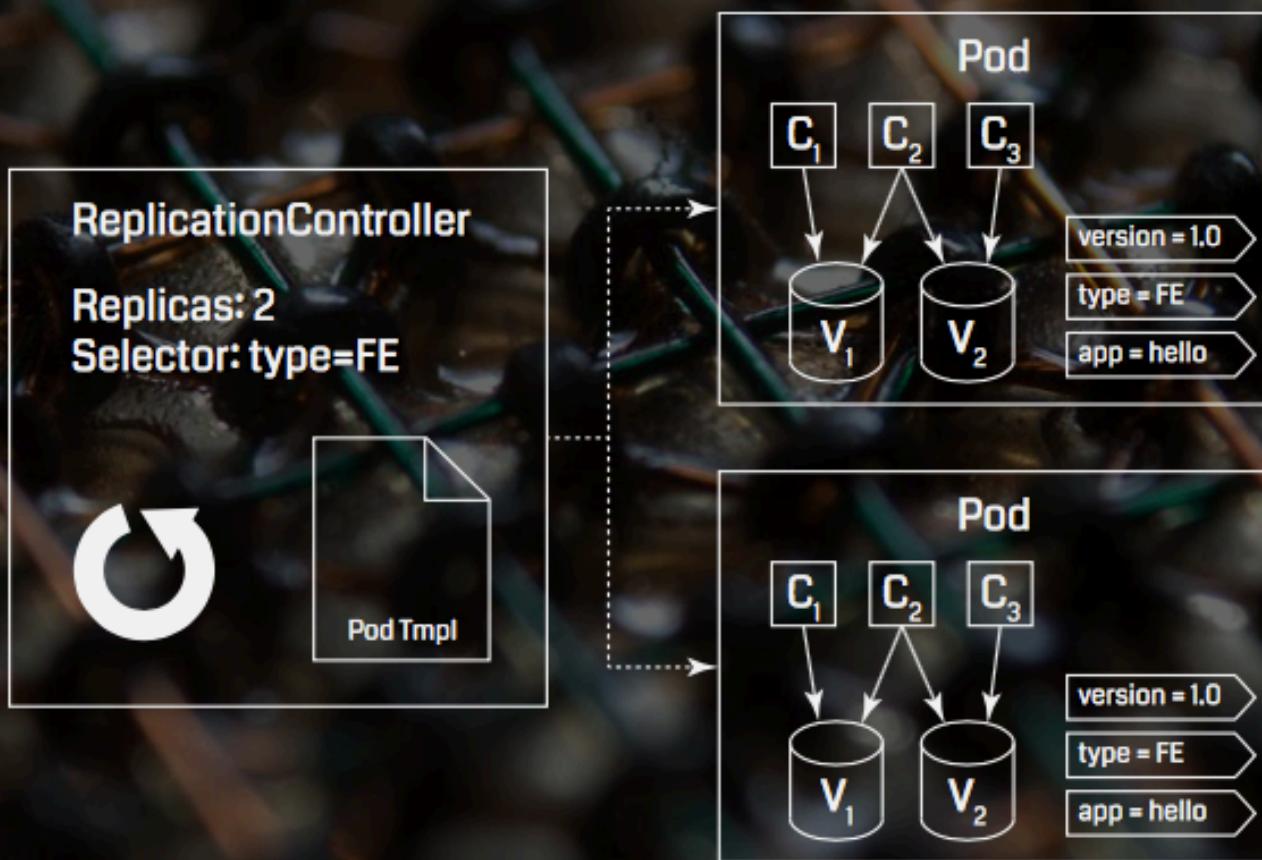


Components: Ambassadors



Core Concepts

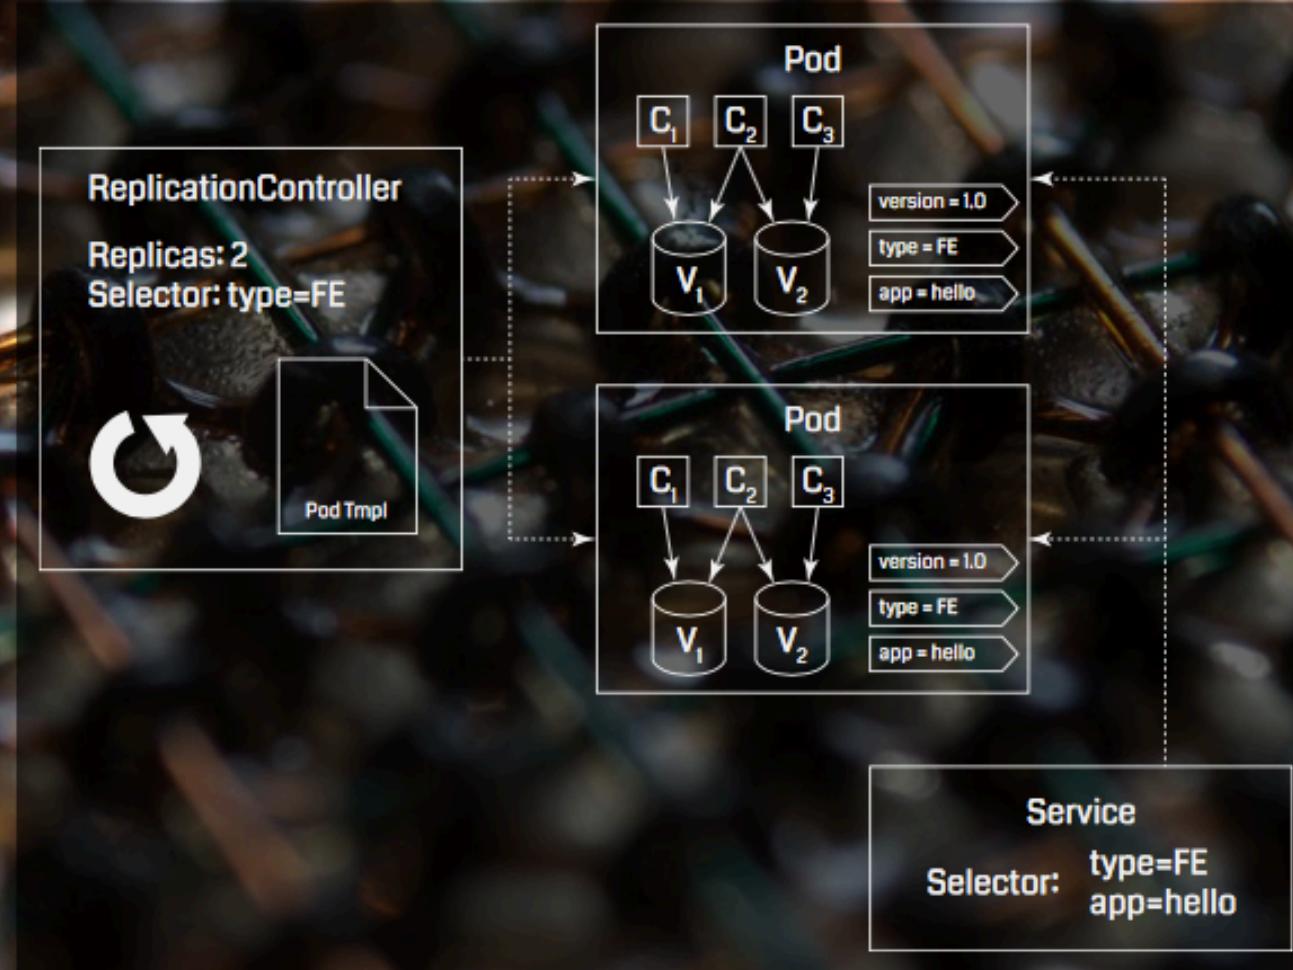
Replication Controller



```
apiVersion: v1
kind: ReplicationController
metadata:
  name: nginx
spec:
  replicas: 3
  selector:
    app: nginx
  template:
    metadata:
      name: nginx
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx
        ports:
        - containerPort: 80
```

Core Concepts

Service



Defining a service

Suppose you have a set of `pods` that each expose port `9376` and carry a label `"app=MyApp"`.

```
{  
    "kind": "Service",  
    "apiVersion": "v1",  
    "metadata": {  
        "name": "my-service"  
    },  
    "spec": {  
        "selector": {  
            "app": "MyApp"  
        },  
        "ports": [  
            {  
                "protocol": "TCP",  
                "port": 80,  
                "targetPort": 9376  
            }  
        ]  
    }  
}
```

In any of these scenarios you can define a `service` without a `selector`:

```
{  
  "kind": "Service",  
  "apiVersion": "v1",  
  "metadata": {  
    "name": "my-service"  
  },  
  "spec": {  
    "ports": [  
      {  
        "protocol": "TCP",  
        "port": 80,  
        "targetPort": 9376  
      }  
    ]  
  }  
}  
  
{  
  "kind": "Endpoints",  
  "apiVersion": "v1",  
  "metadata": {  
    "name": "my-service"  
  },  
  "subsets": [  
    {  
      "addresses": [  
        { "ip": "1.2.3.4" }  
      ],  
      "ports": [  
        { "port": 9376 }  
      ]  
    }  
  ]  
}
```

Discovering services

Environment variables

```
REDIS_MASTER_SERVICE_HOST=10.0.0.11
REDIS_MASTER_SERVICE_PORT=6379
REDIS_MASTER_PORT=tcp://10.0.0.11:6379
REDIS_MASTER_PORT_6379_TCP=tcp://10.0.0.11:6379
REDIS_MASTER_PORT_6379_TCP_PROTO=tcp
REDIS_MASTER_PORT_6379_TCP_PORT=6379
REDIS_MASTER_PORT_6379_TCP_ADDR=10.0.0.11
```

DNS

my-service.my-namespace

Publishing services - service types

Service could be exposed onto an external or internal network. Kubernetes `ServiceTypes` allow you to specify what kind of service you want.

Valid values for the `ServiceType` field are:

- **ClusterIP**: use a cluster-internal IP only - this is the default and is discussed above. Choosing this value means that you want this service to be reachable only from inside of the cluster.
- **NodePort**: on top of having a cluster-internal IP, expose the service on a port on each node of the cluster (the same port on each node). You'll be able to contact the service on any `NodeIP:NodePort` address.
- **LoadBalancer**: on top of having a cluster-internal IP and exposing service on a NodePort also, ask the cloud provider for a load balancer which forwards to the `Service` exposed as a `NodeIP:NodePort` for each Node.

Type LoadBalancer

```
{  
    "kind": "Service",  
    "apiVersion": "v1",  
    "metadata": {  
        "name": "my-service"  
    },  
    "spec": {  
        "selector": {  
            "app": "MyApp"  
        },  
        "ports": [  
            {  
                "protocol": "TCP",  
                "port": 80,  
                "targetPort": 9376,  
                "nodePort": 30061  
            }  
        ],  
        "clusterIP": "10.0.171.239",  
        "loadBalancerIP": "78.11.24.19",  
        "type": "LoadBalancer"  
    },  
    "status": {  
        "loadBalancer": {  
            "ingress": [  
                {  
                    "ip": "146.148.47.155"  
                }  
            ]  
        }  
    }  
}
```

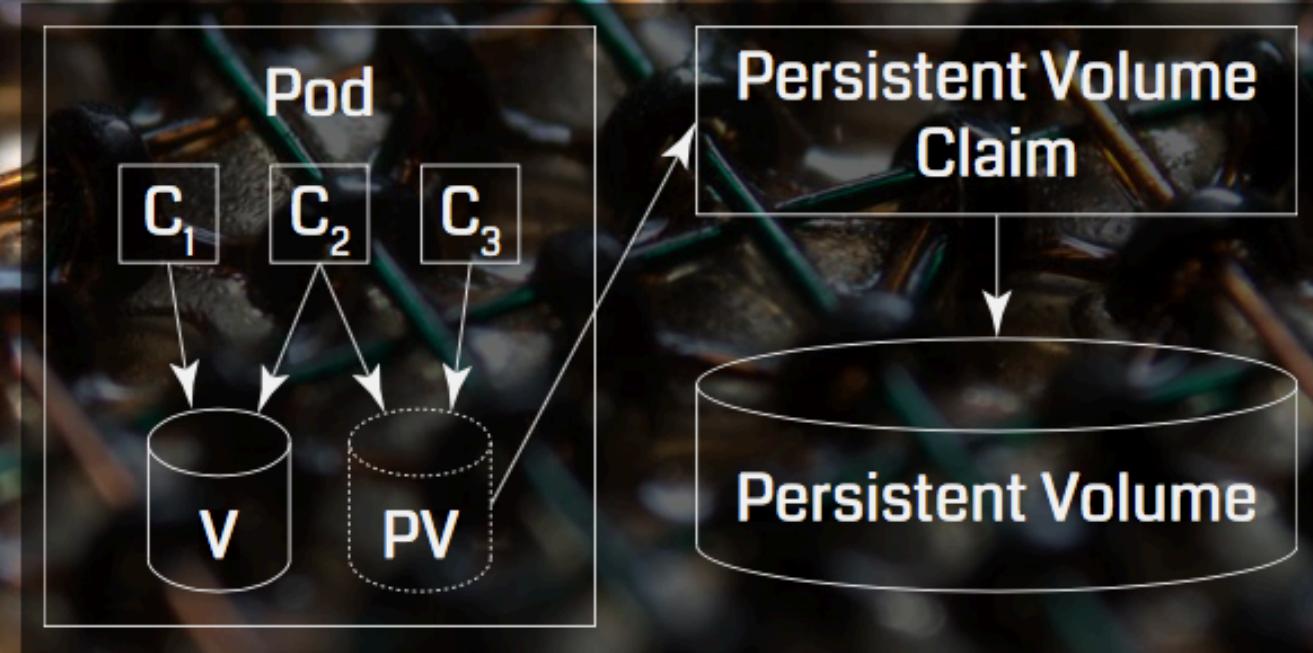
External IPs

If there are external IPs that route to one or more cluster nodes, Kubernetes services can be exposed on those `externalIPs`. `externalIPs` are not managed by Kubernetes and are the responsibility of the cluster administrator. In the example below, `my-service` can be accessed by clients on `80.11.12.10:80` (`externalIP:port`)

```
{  
    "kind": "Service",  
    "apiVersion": "v1",  
    "metadata": {  
        "name": "my-service"  
    },  
    "spec": {  
        "selector": {  
            "app": "MyApp"  
        },  
        "ports": [  
            {  
                "name": "http",  
                "protocol": "TCP",  
                "port": 80,  
                "targetPort": 9376  
            }  
        ],  
        "externalIPs" : [  
            "80.11.12.10"  
        ]  
    }  
}
```

Core Concepts

Persistent Volumes



Types of Persistent Volumes

PersistentVolume types are implemented as plugins. Kubernetes currently supports the following plugins:

- GCEPersistentDisk
- AWSElasticBlockStore
- NFS
- iSCSI
- RBD (Ceph Block Device)
- Glusterfs
- HostPath (single node testing only – local storage is not supported in any way and WILL NOT WORK in a multi-node cluster)

awsElasticBlockStore

An **awsElasticBlockStore** volume mounts an Amazon Web Services (AWS) [EBS Volume](#) into your pod. Unlike **emptyDir**, which is erased when a Pod is removed, the contents of an EBS volume are preserved and the volume is merely unmounted. This means that an EBS volume can be pre-populated with data, and that data can be "handed off" between pods.

Important: You must create an EBS volume using [aws ec2 create-volume](#) or the AWS API before you can use it

There are some restrictions when using an awsElasticBlockStore volume:

- the nodes on which pods are running must be AWS EC2 instances
- those instances need to be in the same region and availability-zone as the EBS volume
- EBS only supports a single EC2 instance mounting a volume

AWS EBS Example configuration

```
apiVersion: v1
kind: Pod
metadata:
  name: test-ebs
spec:
  containers:
    - image: gcr.io/google_containers/test-webserver
      name: test-container
      volumeMounts:
        - mountPath: /test-ebs
          name: test-volume
  volumes:
    - name: test-volume
      # This AWS EBS volume must already exist.
      awsElasticBlockStore:
        volumeID: <volume-id>
        fsType: ext4
```